

# Paper Review - Auto-FuzzyJoin: Auto-Program Fuzzy Similarity Joins Without Labelled Examples

Peng Li, Xiang Cheng, Xu Chu, Yeye He, Surajit Chaudhuri

Carmel Gafa




April 2, 2025

# Fuzzy-join (or similarity join)



- Fuzzy join takes two tables as inputs and identify record pairs that refer to the same entity
- As an example, l1 and r2 refer to the same person
- The concept can be extended to rows consisting of multiple columns

# Fuzzy-join configuration



- Fuzzy-join has been integrated many commercial applications
- These systems are normally not easy to use as they have a big number of configuration parameters
- The extension in Microsoft Excel has 19 options that span across 3 dialog boxes
  - 11 are binary, thus resulting in 2048 possible configuration scenarios
  - 8 continuous, such as thresholds and biases
- In order to execute quality Fuzzy-joins, these configurations must be set up properly by the user

# Theoretical foundation: fuzzy join

Given a **reference table**  $L$  and a table  $R$  containing records that may be **imprecise** or noisy, a **fuzzy join**  $J$  establishes approximate matches between them.

- $J$  connects elements of  $R$  to similar elements in  $L$  based on a chosen **similarity measure** (e.g., Levenshtein distance, cosine similarity, Jaccard similarity).
- Each record  $r \in R$  is mapped to at most one record  $l \in L$ , or **no match at all** (denoted by  $\perp$ ).
- The join is **many-to-one** because multiple records in  $R$  can be associated with the **same** record in  $L$ , but each  $r \in R$  has only **one** possible match.

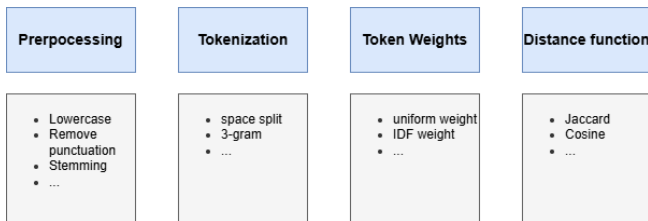
Formally:

$$J : R \rightarrow L \cup \perp$$

# Theoretical foundation: fuzzy join configuration space

A **fuzzy join**  $f$  compares two strings,  $r$  and  $l$ , by computing a distance score that reflects their similarity. The computation of this score is governed by a variety of parameters, forming a **parameter space**.

**Definition:** Each unique combination of these parameters defines a specific **join function**  $f \in \mathcal{F}$ , where  $\mathcal{F}$  is the space of all possible join functions.



# Example: Fuzzy Join Distance Score Computation

**Join Function:**  $f = (L, SP, EW, JD)$

- **L:** Lower-casing (Preprocessing)
- **SP:** Space Tokenization
- **EW:** Equal Weights
- **JD:** Jaccard Distance

**Inputs:**

- $l = \text{"2012 tigers lsu baseball team"}$
- $r = \text{"2012 lsu baseball team"}$

**Tokenization (SP):**

- $l \rightarrow \{2012, tigers, lsu, baseball, team\}$
- $r \rightarrow \{2012, lsu, baseball, team\}$

**Jaccard Distance:**

- $A \cap B = \{2012, lsu, baseball, team\} \rightarrow |A \cap B| = 4$
- $A \cup B = \{2012, tigers, lsu, baseball, team\} \rightarrow |A \cup B| = 5$
- Jaccard Similarity  $= \frac{4}{5} = 0.8$
- Jaccard Distance  $= 1 - 0.8 = 0.2$

**Result:**  $f(l, r) = 0.2$

# Theoretical foundation: threshold and join configuration

- Once the distance  $f(l, r)$  is computed:
  - It is compared to a threshold **compared to a threshold**  $\theta$  to decide whether to join the string pair  $l$  and  $r$ .
  - If  $f(l, r) \leq \theta$ , the pair is considered a **match**.
- Together, the function  $f$  and the threshold  $\theta$  define what the authors call a **join configuration**:

$$C = (f, \theta)$$

- This configuration encapsulates both:
  - How distance is computed.
  - When two strings are considered similar enough to be joined.

**Definition:** A join configuration  $C$  is a 2-tuple  $C = \langle f, \theta \rangle$ , where  $f \in \mathcal{F}$  is a join function, and  $\theta$  is a threshold.

We use  $\mathcal{S} = \{ \langle f, \theta \rangle \mid f \in \mathcal{F}, \theta \in \mathbb{R} \}$  to denote the space of join configurations.

# Theoretical foundation: fuzzy join mapping

Given two tables  $L$  and  $R$ , a join configuration  $C \in \mathcal{S}$  induces a **fuzzy join mapping**  $J_C$ , defined as:

$$J_C(r) = \arg \min_{l \in L, f(l,r) \leq \theta} f(l,r), \forall r \in R$$

That is

- For each record  $r \in R$ , find  $l \in L$  that minimizes the distance  $f(l,r)$ , **only if** that distance is less than or equal to the threshold  $\theta$ .
- If no such  $l \in L$  exists such that  $f(l,r) \leq \theta$ , then  $J_C(r)$  is **empty** — i.e., no match for that record.



# Theoretical foundation: the problem with single join configurations

Real-world data can exhibit **multiple types of variations simultaneously**, such as:

- **Typos**
- **Missing tokens**
- **Extraneous information**

As a result, relying on a **single join configuration** often fails to capture all valid matches, particularly when high **recall** is required.

To handle this diversity, the algorithm uses a **set of join configurations**:

$$U = \{C_1, C_2, \dots, C_K\}$$

Instead of relying on a single configuration, the system computes join results from each one.

This approach allows the system to:

- Accommodate diverse types of variations.
- Improve overall recall by **combining multiple perspectives** on similarity.

L-id	L-Table (Reference Table)		R-id	R-Table (Input Table)
$l_1$	2008 LSU Tigers baseball team	↔	$r_1$	2008 LSU baseball team
$l_2$	2008 LSU Tigers football team	↔	$r_2$	2008 LSU football team
$l_3$	2008 Mississippi State Bulldogs baseball team	↔	$r_3$	2008 Mississippi State <b>Bulldog</b> baseball team
$l_4$	2008 Mississippi State Bulldogs football team	↔	$r_4$	2008 Mississippi State <b>Bulldog</b> football team
$l_5$	...		$r_5$	...
$l_6$	2007 LSU Tigers football team	✗	$r_6$	2007 LSU Tigers baseball team
$l_7$	2007 Wisconsin Badgers football team	✗	$r_7$	2008 Wisconsin Badgers football team
$l_8$	2011 LSU Tigers football team	✗	$r_8$	2010 LSU Tigers football team
$l_9$	2007 LSU Tigers baseball team	✗	$r_9$	2007 LSU Tigers football team

- A **Jaccard distance** with threshold 0.2 works well for pairs like  $(l_1, r_1)$ , which differ by only one or two tokens.
- However, for pairs like  $(l_3, r_3)$  with **spelling variations**, Jaccard similarity is not enough:
  - Jaccard distance  $\approx 0.5 \rightarrow$  too high to match under the 0.2 threshold
  - A more suitable metric is **Edit Distance**, which can better align such pairs.

# Fuzzy Join via Multiple Configurations

- To handle this diversity, the algorithm uses a **set of join configurations**:

$$U = \{C_1, C_2, \dots, C_K\}$$

- Instead of relying on a single configuration, the system computes join results from each.
- This approach allows the system to:
  - Accommodate diverse types of variations.
  - Improve overall recall by **combining multiple perspectives** on similarity.

Given  $L$  and  $R$ , a set of join configurations  $U = \{C_1, C_2, \dots, C_K\}$  induces a **fuzzy join mapping**  $J_U$ , defined as:

$$J_U(r) = \bigcup_{C_i \in U} J_{C_i}(r), \quad \forall r \in R \quad (2)$$

This means that the overall result of the fuzzy join using configuration set  $U$  is the **union** of results from all individual configurations  $C_i \in U$ .

Each configuration  $C_i \in U$  is designed to capture a **specific type of string variation** (e.g., typos, missing tokens, extra tokens).

Two records are considered **joined by the set  $U$**  if and only if they are joined by **at least one** configuration  $C_i \in U$ .

- Each configuration contributes **high-quality joins** targeted at particular data challenges.
- The overall join is more **robust and comprehensive**.