

Paper Review - Auto-FuzzyJoin: Auto-Program Fuzzy Similarity Joins Without Labelled Examples

Peng Li, Xiang Cheng, Xu Chu, Yeye He, Surajit Chaudhuri

Carmel Gafa




April 2, 2025

Fuzzy-join (or similarity join)



- Fuzzy join takes two tables as inputs and identify record pairs that refer to the same entity
- As an example, l1 and r2 refer to the same person
- The concept can be extended to rows consisting of multiple columns

Fuzzy-join configuration



- Fuzzy-join has been integrated many commercial applications
- These systems are normally not easy to use as they have a big number of configuration parameters
- The extension in Microsoft Excel has 19 options that span across 3 dialog boxes
 - 11 are binary, thus resulting in 2048 possible configuration scenarios
 - 8 continuous, such as thresholds and biases
- In order to execute quality Fuzzy-joins, these configurations must be set up properly by the user

Theoretical foundation: fuzzy join

Given a **reference table** L and a table R containing records that may be **imprecise** or noisy, a **fuzzy join** J establishes approximate matches between them.

- J connects elements of R to similar elements in L based on a chosen **similarity measure** (e.g., Levenshtein distance, cosine similarity, Jaccard similarity).
- Each record $r \in R$ is mapped to at most one record $l \in L$, or **no match at all** (denoted by \perp).
- The join is **many-to-one** because multiple records in R can be associated with the **same** record in L , but each $r \in R$ has only **one** possible match.


Formally:

$$J : R \rightarrow L \cup \perp$$

Theoretical foundation: fuzzy join configuration space

A **fuzzy join** f compares two strings, r and l , by computing a distance score that reflects their similarity. The computation of this score is governed by a variety of parameters, forming a **parameter space**.

Definition: Each unique combination of these parameters defines a specific **join function** $f \in \mathcal{F}$, where \mathcal{F} is the space of all possible join functions.



Example: Fuzzy Join Distance Score Computation

Join Function: $f = (L, SP, EW, JD)$

- **L:** Lower-casing (Preprocessing)
- **SP:** Space Tokenization
- **EW:** Equal Weights
- **JD:** Jaccard Distance

Inputs:

- $l = \text{"2012 tigers lsu baseball team"}$
- $r = \text{"2012 lsu baseball team"}$

Tokenization (SP):

- $l \rightarrow \{2012, tigers, lsu, baseball, team\}$
- $r \rightarrow \{2012, lsu, baseball, team\}$

Jaccard Distance:

- $A \cap B = \{2012, lsu, baseball, team\} \rightarrow |A \cap B| = 4$
- $A \cup B = \{2012, tigers, lsu, baseball, team\} \rightarrow |A \cup B| = 5$
- Jaccard Similarity $= \frac{4}{5} = 0.8$
- Jaccard Distance $= 1 - 0.8 = 0.2$

Result: $f(l, r) = 0.2$

Theoretical foundation: threshold and join configuration

- Once the distance $f(l, r)$ is computed:
 - It is compared to a threshold **compared to a threshold** θ to decide whether to join the string pair l and r .
 - If $f(l, r) \leq \theta$, the pair is considered a **match**.
- Together, the function f and the threshold θ define what the authors call a **join configuration**:

$$C = (f, \theta)$$

- This configuration encapsulates both:
 - How distance is computed.
 - When two strings are considered similar enough to be joined.

Definition: A join configuration C is a 2-tuple $C = \langle f, \theta \rangle$, where $f \in \mathcal{F}$ is a join function, and θ is a threshold.

We use $\mathcal{S} = \{ \langle f, \theta \rangle \mid f \in \mathcal{F}, \theta \in \mathbb{R} \}$ to denote the space of join configurations.

Theoretical foundation: fuzzy join mapping

Given two tables L and R , a join configuration $C \in \mathcal{S}$ induces a **fuzzy join mapping** J_C , defined as:

$$J_C(r) = \arg \min_{l \in L, f(l,r) \leq \theta} f(l,r), \forall r \in R$$

That is

- For each record $r \in R$, find $l \in L$ that minimizes the distance $f(l,r)$, **only if** that distance is less than or equal to the threshold θ .
- If no such $l \in L$ exists such that $f(l,r) \leq \theta$, then $J_C(r)$ is **empty** — i.e., no match for that record.

Theoretical foundation: the problem with single join configurations

Real-world data can exhibit **multiple types of variations simultaneously**, such as:

- **Typos**
- **Missing tokens**
- **Extraneous information**

As a result, relying on a **single join configuration** often fails to capture all valid matches, particularly when high **recall** is required.

To handle this diversity, the algorithm uses a **set of join configurations**:

$$U = \{C_1, C_2, \dots, C_K\}$$

Instead of relying on a single configuration, the system computes join results from each one.

This approach allows the system to:

- Accommodate diverse types of variations.
- Improve overall recall by **combining multiple perspectives** on similarity.

L-id	L-Table (Reference Table)		R-id	R-Table (Input Table)
l_1	2008 LSU Tigers baseball team	↔	r_1	2008 LSU baseball team
l_2	2008 LSU Tigers football team	↔	r_2	2008 LSU football team
l_3	2008 Mississippi State Bulldogs baseball team	↔	r_3	2008 Mississippi State Bulldog baseball team
l_4	2008 Mississippi State Bulldogs football team	↔	r_4	2008 Mississippi State Bulldog football team
l_5	...		r_5	...
l_6	2007 LSU Tigers football team	↔	r_6	2007 LSU Tigers baseball team
l_7	2007 Wisconsin Badgers football team	↔	r_7	2008 Wisconsin Badgers football team
l_8	2011 LSU Tigers football team	↔	r_8	2010 LSU Tigers football team
l_9	2007 LSU Tigers baseball team	↔	r_9	2007 LSU Tigers football team

- A **Jaccard distance** with threshold 0.2 works well for pairs like (l_1, r_1) , which differ by only one or two tokens.
- However, for pairs like (l_3, r_3) with **spelling variations**, Jaccard similarity is not enough:
 - Jaccard distance $\approx 0.5 \rightarrow$ too high to match under the 0.2 threshold
 - A more suitable metric is **Edit Distance**, which can better align such pairs.

Theoretical foundation: fuzzy join via multiple configurations

- To handle this diversity, the algorithm uses a **set of join configurations**:

$$U = \{C_1, C_2, \dots, C_K\}$$

- Instead of relying on a single configuration, the system computes join results from each.
- This approach allows the system to:
 - Accommodate diverse types of variations.
 - Improve overall recall by **combining multiple perspectives** on similarity.

Given L and R , a set of join configurations $U = \{C_1, C_2, \dots, C_K\}$ induces a **fuzzy join mapping** J_U , defined as:

$$J_U(r) = \bigcup_{C_i \in U} J_{C_i}(r), \quad \forall r \in R \quad (2)$$

This means that the overall result of the fuzzy join using configuration set U is the **union** of results from all individual configurations $C_i \in U$.

Each configuration $C_i \in U$ is designed to capture a **specific type of string variation** (e.g., typos, missing tokens, extra tokens).

Two records are considered **joined by the set U if and only if** they are joined by **at least one** configuration $C_i \in U$.

- Each configuration contributes **high-quality joins** targeted at particular data challenges.
- The overall join is more **robust and comprehensive**.

Evaluating Join Quality: Precision

Given two tables R and L , and a **space of join configurations** \mathcal{S} , the objective is to find a subset $U \subseteq \mathcal{S}$ that produces **good fuzzy join results**.

Let:

- J_U be the fuzzy join mapping induced by configuration set U
- J_G be the **ground truth** join mapping — the ideal join result

Precision measures how many of the predicted joins are correct:

$$\text{precision}(U) = \frac{\underbrace{|\{r \in R \mid J_U(r) \neq \emptyset, J_U(r) = J_G(r)\}|}_{\text{True Positives (TP)}}}{\underbrace{|\{r \in R \mid J_U(r) \neq \emptyset\}|}_{\text{TP + FP (all predicted joins)}}} \quad (3)$$

- **Numerator (TP)**: Records where a join was predicted and it matched the ground truth.
- **Denominator (TP + FP)**: All records where a join was predicted (correct or not).

Evaluating Join Quality: Recall

Recall measures how many of the correct (ground truth) joins were successfully predicted:

$$\text{recall}(U) = \underbrace{|\{r \in R \mid J_U(r) \neq \emptyset, J_U(r) = J_G(r)\}|}_{\text{True Positives (TP)}} \quad (4)$$

- This is the **absolute count of True Positives**, i.e., records for which:
 - A join was predicted ($J_U(r) \neq \emptyset$), and
 - It matches the ground truth ($J_U(r) = J_G(r)$)

False Negatives (FN) — cases where a correct join was missed — are defined as:

$$\text{FN} = |\{r \in R \mid J_G(r) \neq \emptyset, J_U(r) = \emptyset\}|$$

Note: The denominator $TP + FN$ is constant across all U for a fixed dataset, so it is omitted in comparisons.

Estimating Precision/Recall for a Single Join Configuration

Given:

- A **single join configuration** $C = \langle f, \theta \rangle$
- Two tables:
 - L : reference table
 - R : query table

Assumption: Complete Reference Table L

- L is assumed to contain **all possible true matches** for records in R .
- Ensures that for each $r \in R$, there exists a correct match $l \in L$.
- Simplifies analysis by reducing the chance of missing true positives due to an incomplete reference.

Geometric View of the Distance Function f

- Join function f embeds records into a **metric space**.
- Records are modeled as points on a **unit grid**.
- Each $l \in L$ is surrounded by **close variants** (differing by a token, character, etc.).

Analogy: Stars and Planets

- Reference records $l \in L$ are like **stars** on a grid.
- Query records $r \in R$ are like **planets** that orbit these stars.
- Identifying the best join $J_C(r)$ is like determining **which star a planet orbits**.

Safe Joins and the Geometry of Fuzzy Matching

Safe Joins with a Complete L

- Define the **grid width** w : typical distance between a record l and its closest neighbors in L .
- A join is considered **safe** if the distance $d = f(l, r)$ satisfies:

$$d < \frac{w}{2}$$

- This guarantees that r lies **closer to its true match** l than to any other reference point.

Why This Matters:

- Ensures high **precision** — avoiding false positives caused by ambiguous joins.
- Avoids joining r to an incorrect l' that lies at a similar distance.

Analogy:

- A planet that lies **equidistant** between two stars (at $\frac{w}{2}$ each) **cannot be confidently claimed by either**.
- In fuzzy joining, such cases are inherently **ambiguous** and risky to resolve.